Team tasks and interfaces

One mode, open mouth, real time. Austin calls a judge through Twilio and mouths the whole conversation; the camera, the EMG model and Claude read it together every 100 ms; live captions show both sides; Cartesia speaks Austin's words into the call in his cloned voice, with his facial emotion, while he is still mouthing. Everything runs on Austin's laptop (RTX 5060, 115 W), with the lowest delay we can get at every step.

How the whole system works: REALTIME_PIPELINE.md. The EMG model in detail: SVR1_SPEC.md.

The demo

The pitch is separate: Austin and Frank give it with their real voices. The product demo is a phone call.

1. Austin calls the judge, who is standing in front of us, through Twilio. The judge speaks normally into their phone; their words appear in the captions.
2. Austin answers only by mouthing: mouth open, no sound. His words appear on screen as he mouths, grey then black, and his cloned voice speaks them into the call while he is still going.
3. His facial expression changes the voice's emotion.
4. The dashboard shows the real latency, which reader settled each word, and what Claude changed.
5. If the system can't tell what he said, the screen says "Please repeat" rather than guessing.

Nothing is trained on the pitch or on a script: the judge can ask anything.

Who does what

Person	Owns
Evan	Camera reader: real-time Auto-AVSR every 100 ms, mouthing activity and segments, facial emotion (EVAN_REALTIME_LIP_READING.md)
Austin	SVR-1 training, live EMG reader, fusion decoder and confidence locking, logging and replay, integration
Frank	Cartesia (Austin's voice clone, Sonic streaming voice, Ink for the judge), the Twilio call, the captions dashboard
Thai	EMG rig (4 channels, sites 3, 2, 6, 1) and its live filter, the Claude service (predict + repair)

Austin

1. Retarget SVR-1 to the rig's 4 sites (input slots 0, 1, 2, 5, the existing `4ch` channel set). Drop the closed-mouth parts: the fixed-phrase head and closed-lips transfer.

2. Switch all EMG preprocessing to Thai's causal filter (`rig/filters.py`), Berkeley pre-training data included: re-cache Berkeley through it, in order, instead of Gaddy's `filtfilt`.
3. Add the speaking-or-not output and partial-sentence training.
4. Train on the cluster: one full SVR-1 run plus a Gaddy-architecture baseline on the same 4 sites.
5. The recording tool: camera + EMG together with the session prompts (see Recording sessions).
6. Live EMG reader: subscribe to `emg.raw`, filter, run SVR-1 every 100 ms, compute pace and effort.
7. Fusion decoder: the joint search over camera and EMG probabilities, with look-alike expansion, Claude's predicted words and the gate.
8. Confidence locking: compute the lock ingredients per word, train the lock model on sessions 1 and 2, and set the dynamic threshold from Claude's predictability.
9. "Please repeat" when the final pass still isn't confident, with the reasons in the debug log.
10. The logger, which subscribes to every topic and writes everything to disk, and a replay tool.
11. Fine-tune SVR-1 on sessions 1 and 2, and tune the decoder, gate and lock model; test everything on session 3.

Evan

Follow EVAN_REALTIME_LIP_READING.md. You can prototype on the MacBook (Lipflow was built for Macs), but timing and real tests run only on Austin's laptop.

Frank

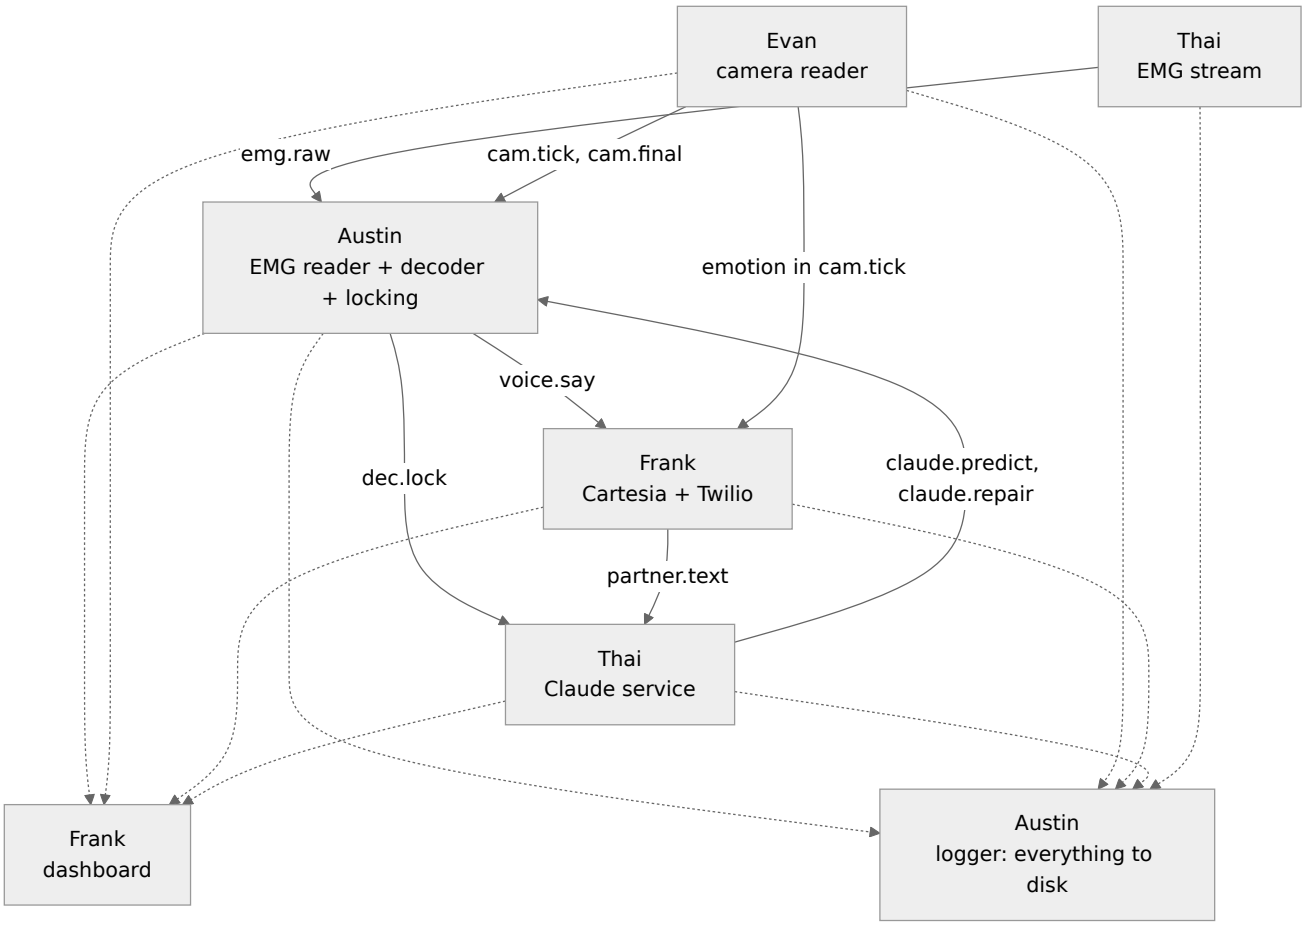
1. Instant Voice Clone of Austin, from his own recordings, with his consent.
2. Sonic 3.6 over the WebSocket: one context per sentence, words streamed in with "more to come" as they arrive on `voice.say`, `max_buffer_delay_ms` kept low. Speed, volume and emotion go through `generation_config`.
3. Twilio:
    - Austin's call to the judge, with a two-way audio stream to the laptop.
    - Cartesia's audio goes into the call; ask Cartesia for phone-format audio, or convert it.
    - The judge's incoming audio goes to Cartesia Ink, which publishes `partner.text`.
4. The captions dashboard:
    - captions for both speakers, with Austin's words grey until locked, then black;
    - Claude's edits highlighted, and a raw-vs-fixed panel;
    - a live latency meter;
    - "Please repeat" with the reason;
    - the EMG traces.

Thai

1. Rig with 4 boards on A0 to A3 (sites 3, 2, 6, 1) and one REF per board, as in HARDWARE.md.
2. Firmware, reader and live filter (`rig/`): 1 kHz, raw frames with a counter, causal hum filter. They move from 6 to 4 channels with the recording tool.
3. Live: the reader stamps samples with `time.monotonic()` and publishes `emg.raw` every 10 ms.
4. Safety: the laptop runs on battery whenever the pads are on skin; never charge with the Uno connected.
5. The Claude service:
    - **Predict:** on every final `partner.text` and every ~2 s while Austin mouths, publish `claude.predict`: likely words and phrases, plus how predictable the next words are.
    - **Repair:** on each `dec.lock`, publish `claude.repair`.
    - **How it calls Claude:** `claude-opus-5-5` at low effort, structured outputs, prompt caching (rules and conversation kept as an append-only prefix).
    - **Prompt rules:** swap a word only for one of the alternatives given or a lip look-alike, never add ideas, answer "please repeat" when unsure.

Interfaces

The parts run as separate small programs on the laptop and talk through ZeroMQ publish/subscribe over local sockets (`ipc://`), with messages encoded as msgpack (numpy arrays as raw bytes). A message takes well under a millisecond. All timestamps are `time.monotonic()` seconds on the same laptop.



Topic	From → to	Rate	Main fields
`emg.raw`	Thai → Austin	every 10 ms	`t0`, `fs` = 1000, `samples` (n × 4, raw 0 to 1023), `counter`
`cam.tick`	Evan → Austin, Frank	every 100 ms	`seg_id`, `seg_start`, `frame_t` (T), `topk_ids` (T × 32), `topk_logp` (T × 32), `blank_logp` (T), `greedy`, `activity` (T), `face_ok`, `emotion` {`label`, `confidence`, `scores`}, `timing_ms`
`cam.final`	Evan → Austin	at segment end	`seg_id`, `nbest` [{`text`, `score`}]
`dec.hyp`	Austin → dashboard	every 100 ms	`utt_id`, `locked` [words], `unlocked` [words], `confidence` per word, `source` per word (camera / EMG / Claude)
`dec.lock`	Austin → Claude, dashboard	on lock	`utt_id`, `words` [{`w`, `conf`, `alts`, `cam`, `emg`, `t`}]
`dec.repeat`	Austin → dashboard	at sentence end	`utt_id`, `reason` (which reader disagreed, confidences)
`claude.predict`	Claude → Austin	on judge line, ~2 s	`words` [{`w`, `weight`}], `phrases`, `predictability` (0 to 1), `ttl_s`
`claude.repair`	Claude → Austin, dashboard	on reply	`utt_id`, `edits` [{`index`, `from`, `to`}], `please_repeat`

Topic	From → to	Rate	Main fields
`voice.say`	Austin → Frank	words leaving the repair window	`utt_id`, `words`, `final`, `pace`, `effort`
`partner.text`	Frank → Claude, dashboard	streaming	`text`, `is_final`, `t`
`timing.*`	everyone → logger	every event	the stage and its duration in ms

Austin applies Claude's repairs before forwarding words on `voice.say`, so the voice never needs to know about Claude. The logger subscribes to every topic.

Order of work

The system must work end to end early and get better, not work only at the very end.

Block	Goal	Evan	Austin	Frank	Thai
1	Record + skeleton	E0, E1, E7 recording	Session 1; bus + logger; SVR-1 retarget; cluster runs start	Voice clone; Sonic speaks typed text	Rig for session 1; `emg.raw` live
2	Camera → voice, end to end	E2, E4	Decoder on camera alone; locking v1 → `voice.say`	Streaming words in; Twilio call + Ink	Claude predict on real conversation
3	Session 2 + EMG joins	E3	Session 2; fine-tune SVR-1; EMG into the decoder	Dashboard	Claude repair
4	Learn and tune	E5, E6	Train the lock model and gate on sessions 1 and 2	Emotion, pace, volume	Prompt tuning on logs
5	Test and rehearse	Replay tests	Session 3 in demo conditions; WER and latency on it	Full calls with a teammate as judge	Demo-room check of light and noise

Fallback at every point: camera + Claude + voice works without EMG. If EMG isn't ready by the demo, that is what we show, and nothing else changes.

Recording sessions

Procedure: RECORDING.md; sentences: recording/. Sessions are recorded with the recording tool (`recorder/`), which records the camera and Thai's EMG stream on one clock and steps through the prompts. All varied conversational sentences, mouth open, no sound. The pads come off and go back on between sessions.

Session	Sentences	Use
1 (now)	250, the last 50 faster than natural	Training

Session	Sentences	Use
2	250, the last 50 faster than natural	Training
3 (last, in demo conditions)	100, at demo pace	Test only, never trained on

Each session also records 10 rest clips of about 3 s (still face, swallow, nod, smile, and so on), which teach the model what is not speech. Sessions 1 and 2 take about 30 minutes each, session 3 about 12. A minute of calibration before the demo refits SVR-1's adapter.

Training

1. **Pretraining (cluster):** SVR-1 on Berkeley's data, re-filtered with `rig/filters.py` and retargeted to the rig's 4 sites, with the Gaddy-architecture baseline on the same 4 sites as the fallback and the comparison. Scripts and runbook are in `calhacks2026/cluster/`.
2. **Fine-tuning:** SVR-1 on sessions 1 and 2. It is small enough for the 5090 or the cluster.
3. **The decoder's weights, the gate and the lock model:** fitted on sessions 1 and 2.
4. **The test:** everything is measured on session 3.
5. **Auto-AVSR and Claude are not trained.**

Numbers to show

These come from Anthropic's track brief:

- Word error rate on session 3 for camera alone, camera + EMG, and camera + EMG + Claude.
- Delay from mouthing a word to seeing it locked, and to hearing it (not counting the phone network).
- How often the system says "Please repeat".
- How often the judge understands on the first try.

Still open

- **Battery life:** the whole system runs on battery during recordings and the demo. Measure how long a full run lasts, and charge between sessions with the Uno unplugged.
- **Phone audio:** Frank owns the Twilio ↔ Cartesia audio formats.